Talha Zafar

Senior Backend Engineer
DevOps & Infrastructure Engineer

CONTACT

- +92 307 357 6146
- talha.zafar.j@gmail.com
- linkedin.com/in/talhazafarj
- github.com/talhazafarj
- Lahore, Pakistan

TECHNICAL SKILLS

Backend & APIs

Python · FastAPI · Django · Flask · NestJS · Node.js · gRPC · REST · GraphQL · asyncio · Go · TypeScript

DevOps & Infra

Kubernetes · K3s · Docker · LXD · GitHub Actions · nginx · cert-manager · Helm · Bash · Linux · iptables · fail2ban

Data & Cloud

PostgreSQL · MongoDB · Redis · AWS · GCP · SQLAlchemy · Beanie · vLLM · V100 GPU

Security & Monitoring

OWASP Top 10 · Vault · ECDH · ABE · HMAC · JWT · Trivy · Gitleaks · Zabbix · SonarQube · Grafana

EDUCATION

B.Sc. Computer Science
FAST-NUCES — Lahore
2018 – 2022 · OS · Networks · Distributed Systems

LANGUAGES

English — C1 Professional
Urdu — Native

AVAILABILITY

- Open to new opportunities
- On-site · Hybrid · Remote
- Notice: 1 month (negotiable)
- talhazafarjutt.github.io

PROFESSIONAL SUMMARY

Backend and infrastructure engineer with **4+ years** production experience in high-throughput distributed systems. Expert in **Python microservices** (FastAPI, Django, Flask) and **Node.js/NestJS**. Sole bare-metal infrastructure owner operating 4× **NVIDIA Tesla V100 GPU cluster**, K3s Kubernetes, and zero-trust secrets pipelines. Deep DevSecOps background: OWASP hardening, HashiCorp Vault, SAST/DAST CI/CD. Currently building AI agent platforms and LLM inference infrastructure at production scale.

PROFESSIONAL EXPERIENCE

DevOps & Infrastructure Engineer
Jan 2026 – Present
Augai · Infrastructure & Security Consulting · Remote

- Sole infra owner for bare-metal Ubuntu 22.04 server: **4× V100 GPUs**, K3s Kubernetes, 8+ LXD containers — maintaining **99.8% uptime**.
- Designed K3s cluster: nginx ingress, cert-manager TLS automation, CoreDNS service discovery, GPU passthrough.
- Deployed HashiCorp Vault (AppRole auth, dynamic secrets) in isolated LXD container — **GDPR data residency** compliant.
- Full security hardening: forensic audit (clean), iptables firewall, fail2ban, kernel lockdown, rsyslog centralised across all containers.
- GitHub Actions CI/CD with SAST (SonarQube), container scanning (Trivy), and secret detection (Gitleaks).

K3s · HashiCorp Vault · LXD · iptables · Zabbix · GitHub Actions · Trivy · vLLM · cert-manager

Backend Engineer
Jan 2024 – Dec 2025
BoolMind · AI & Automation Platform · Remote

- Architected production microservices backend: **10k+ daily requests**, **sub-100ms P95** latency, FastAPI + asyncio.
- gRPC inter-service communication across 6 services — **40% less overhead** vs REST; Redis caching cut DB load by **60%**.
- Deployed LLM inference on 4× V100: vLLM serving Qwen2.5-Coder-32B, FastAPI load-balancer, nginx API gateway.
- Implemented OWASP Top 10, ECDH key exchange, ABE encryption for multi-tenant data isolation.
- Stripe security hardening post-carding attack: Radar rules, Redis rate-limiting, HMAC auth — **zero fraud post-deployment**.

FastAPI · gRPC · Redis · PostgreSQL · OWASP · ECDH · ABE · vLLM · asyncio

Junior Backend Engineer
Sep 2022 – Dec 2023
Techleadz · Software Development Agency · Lahore

- Built and scaled RESTful APIs (Django REST, NestJS) for SaaS products with **50k+ registered users**.
- PostgreSQL multi-tenant schemas with row-level security — **35% query time reduction**; MongoDB Beanie ODM for **500k+ daily events**.
- Integrated Stripe, Firebase, OAuth2 for payments, notifications, and authentication end-to-end.

Django · NestJS · PostgreSQL · MongoDB · Redis · Stripe · OAuth2

KEY PROJECTS

LLM Inference Platform (Codsy)
Production multi-GPU cluster on 4× Tesla V100 with K3s orchestration, vLLM per GPU, nginx API gateway, FastAPI load-balancer, and drop-in Python SDK. Serving live LLM traffic for active development team.

Zero-Trust Secrets Pipeline
HashiCorp Vault in LXD with AppRole auth, dynamic credential rotation for SSH and SonarQube CI/CD. Granular policy-based access — GDPR data residency compliant.

AI Agent Microservices Platform
6-service gRPC backend with Redis job queuing, ECDH + ABE multi-tenant isolation. Sub-100ms P95 latency at 10k+ daily API requests.